

AI for Medicine Course 3 Week 1 lecture notebook - Model Training/Tuning Basics with Sklearn

Welcome to this exercise! You're going to be exploring the `sklearn` library, including an overview of its underlying data types and methods for tweaking a model's hyperparameters. You'll be using the same data from the previous lecture notebook. Let's get started!

Packages

First import all the packages that you need for this assignment.

- `pandas` is what you'll use to manipulate your data
- `numpy` is a library for mathematical and scientific operations
- `sklearn` has many efficient tools for machine learning and statistical modeling
- `itertools` helps with hyperparameter (grid) searching

Import Packages

Run the next cell to import all the necessary packages.

```
In [1]: # Import packages
import pandas as pd
import numpy as np
import itertools

# Set the random seed for consistent output
np.random.seed(18)

# Read in the data
data = pd.read_csv("dummy_data.csv", index_col=0)
```

Train/Test Split

```
In [2]: # Import module to split data
from sklearn.model_selection import train_test_split

# Get the label
y = data.outcome

# Get the features
X = data.drop('outcome', axis=1)

# Get training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
= 0.25)
print(f"Number of observations for training: {y_train.size}")
print(f"Number of observations for testing: {y_test.size}")
```

```
Number of observations for training: 37
Number of observations for testing: 13
```

Model Fit and Prediction

Let's fit a logistic regression to the training data. Sklearn allows you to provide arguments that override the defaults.

The default solver is lbfgs.

- Lbfgs stands for '[Limited Memory BFGS](https://en.wikipedia.org/wiki/Limited-memory_BFGS)' (https://en.wikipedia.org/wiki/Limited-memory_BFGS), and is an efficient and popular method for fitting models.
- The solver is set explicitly here for learning purposes; if you do not set the solver parameter explicitly, the [LogisticRegression](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html) (https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html) function will use its default solver, which is lbfgs as well.

```
In [3]: from sklearn.linear_model import LogisticRegression

lr = LogisticRegression(solver='lbfgs')
lr.fit(X_train, y_train)
```

```
Out[3]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                           intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                           penalty='l2', random_state=None, solver='lbfgs', tol=0.001,
                           verbose=0, warm_start=False)
```

When it fits the training data, `sklearn` also prints out the model's hyperparameters.

- Here, these are the default hyperparameters for `sklearn`'s logistic regression classifier.
- Another way to check these parameters is the `get_params()` method of the classifier.

You should spend some time checking out the [documentation \(https://scikit-learn.org/stable/supervised_learning.html#supervised-learning\)](https://scikit-learn.org/stable/supervised_learning.html#supervised-learning) to get a deeper understanding of what's going on. One important thing to note is that each classifier has different hyperparameters.

Prediction

To predict with the classifier, use the `predict()` method.

- This returns a `numpy` array containing the predicted class for each observation in the test set, as you can see by running the next cell:

```
In [4]: # Use the trained model to predict labels from the features of the
        # test set
        predictions = lr.predict(X_test)

        # View the prediction type, shape, and print out a sample prediction
        print(f"predictions is of type: {type(predictions)}")
        print(f"predictions has shape: {predictions.shape}")
        print(f"predicted class for 10th element in test set: {predictions[9]}")

        predictions is of type: <class 'numpy.ndarray'>
        predictions has shape: (13,)
        predicted class for 10th element in test set: 0
```

Prediction probability

When a model predicts that a label is 1 rather than 0, it may help you to know if the model was predicting 1 with a 51% probability or 90% probability; in other words, how confident is that prediction?

You can get the model's probability of predicting each of the class.

- To do this, use the `predict_proba()` method.
- The resulting array will have a shape that matches the number of classes for the target variable.

```
In [5]: prediction_probs = lr.predict_proba(X_test)
print(f"prediction_probs is of type: {type(prediction_probs)}")
print(f"prediction_probs has shape: {prediction_probs.shape}")
print(f"probabilities for first element in test set: {prediction_probs[0]}")

prediction_probs is of type: <class 'numpy.ndarray'>
prediction_probs has shape: (13, 2)
probabilities for first element in test set: [0.42348297 0.57651703]
```

There are 13 patients in the test set. Each patient's label could be either 0 or 1, so the prediction probability has 13 rows and 2 columns. To know which column refers to label 0 and which refers to label 1, you can check the `.classes_` attribute.

```
In [6]: lr.classes_
```

```
Out[6]: array([0, 1])
```

Since the order of the `classes_` array is 0, then 1, column 0 of the prediction probabilities has label 0, and column 1 has label 1.

Let's print these for the first 5 elements of the dataset:

```
In [7]: for i in range(5):
        print(f"Element number: {i}")
        print(f"Predicted class: {predictions[i]}")
        print(f"Probability of predicting class 0: {prediction_probs[i][0]}")
        print(f"Probability of predicting class 1: {prediction_probs[i][1]}\n")
```

```
Element number: 0
Predicted class: 1
Probability of predicting class 0: 0.42348296737845204
Probability of predicting class 1: 0.576517032621548
```

```
Element number: 1
Predicted class: 1
Probability of predicting class 0: 0.4914968703166631
Probability of predicting class 1: 0.5085031296833369
```

```
Element number: 2
Predicted class: 1
Probability of predicting class 0: 0.483088763245346
Probability of predicting class 1: 0.516911236754654
```

```
Element number: 3
Predicted class: 0
Probability of predicting class 0: 0.86953653498578
Probability of predicting class 1: 0.13046346501422001
```

```
Element number: 4
Predicted class: 0
Probability of predicting class 0: 0.8470774295731573
Probability of predicting class 1: 0.15292257042684265
```

You can see here that the predicted class matches the class with a higher probability of being predicted. Since you're dealing with numpy arrays, you can simply slice them and get specific information, such as the probability of predicting class 1 for all elements in the test set:

```
In [8]: # Retrieve prediction probabilities for label 1, for all patients
        prediction_probs[:, 1]
```

```
Out[8]: array([0.57651703, 0.50850313, 0.51691124, 0.13046347, 0.15292257,
               0.26162479, 0.50831618, 0.3190805 , 0.37250246, 0.47736442,
               0.15743244, 0.51193665, 0.26832495])
```

Tuning the Model

Most of the time, the predictive power of a classifier can be increased if a good set of hyperparameters is defined. This is known as model tuning.

For this process, you'll need a classifier, an appropriate evaluation metric, and a set of parameters to test. Since this is a dummy example, you'll use the default metric for the logistic regression classifier: the **mean accuracy**.

Mean Accuracy

Mean Accuracy is the number of correct predictions divided by total predictions. This can be computed with the `score()` method.

Let's begin by checking the performance of your out-of-the-box logit classifier:

```
In [9]: lr.score(X_test, y_test)
```

```
Out[9]: 0.6153846153846154
```

Let's say you want to tweak this model's default parameters. You can pass a dictionary containing the values you specify to the classifier when you instantiate it. Notice that these must be passed as keyword arguments, or `kwargs`, which are created by using the `**` prefix:

```
In [10]: # Choose hyperparameters and place them as key-value pairs in a dictionary
params = {
    'solver': 'liblinear',
    'fit_intercept': False,
    'penalty': 'l1',
    'max_iter': 500
}

# Pass in the dictionary as keyword arguments to the model
lr_tweaked = LogisticRegression(**params)

# Train the model
lr_tweaked.fit(X_train, y_train)

# View hyper-parameters
print(f"Tweaked hyperparameters: {lr_tweaked.get_params()}\n")

# Evaluate the model with the mean accuracy
print(f"Mean Accuracy: {lr_tweaked.score(X_test, y_test)}")

Tweaked hyperparameters: {'C': 1.0, 'class_weight': None, 'dual': False, 'fit_intercept': False, 'intercept_scaling': 1, 'max_iter': 500, 'multi_class': 'ovr', 'n_jobs': 1, 'penalty': 'l1', 'random_state': None, 'solver': 'liblinear', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}

Mean Accuracy: 0.5384615384615384
```

The model with the tweaked parameters is worse than the original! However, there might still be some combination of parameters that increase the predictive power of your logit classifier.

Try different hyperparameters

Testing this can be daunting considering all the possible parameter combinations. Let's try something

To get started, you'll apply `itertools.product()` to create all the combinations of parameters.

- Notice that the iterable (in this case a list of the lists of parameters) must be passed as `*args` to the `product()` function.

```
In [11]: # Choose hyperparameters and place in a dictionary
hyperparams = {
    'solver': ["liblinear"],
    'fit_intercept': [True, False],
    'penalty': ["l1", "l2"],
    'class_weight': [None, "balanced"]
}
# Get the values of hyperparams and convert them to a list of lists
hp_values = list(hyperparams.values())
hp_values
```

```
Out[11]: [['liblinear'], [True, False], ['l1', 'l2'], [None, 'balanced']]
```

```
In [12]: # Get every combination of the hyperparameters
for hp in itertools.product(*hp_values):
    print(hp)

('liblinear', True, 'l1', None)
('liblinear', True, 'l1', 'balanced')
('liblinear', True, 'l2', None)
('liblinear', True, 'l2', 'balanced')
('liblinear', False, 'l1', None)
('liblinear', False, 'l1', 'balanced')
('liblinear', False, 'l2', None)
('liblinear', False, 'l2', 'balanced')
```



```
In [13]: # Loop through the combinations of hyperparams
for hp in itertools.product(*hp_values):

    # Create the model with the hyperparams
    estimator = LogisticRegression(solver=hp[0],
                                   fit_intercept=hp[1],
                                   penalty=hp[2],
                                   class_weight=hp[3])

    # Fit the model
    estimator.fit(X_train, y_train)
    print(f"Parameters used: {hp}")
    print(f"Mean accuracy of the model: {estimator.score(X_test, y_
test)}\n")
```

```
Parameters used: ('liblinear', True, 'l1', None)
Mean accuracy of the model: 0.5384615384615384
```

```
Parameters used: ('liblinear', True, 'l1', 'balanced')
Mean accuracy of the model: 0.46153846153846156
```

```
Parameters used: ('liblinear', True, 'l2', None)
Mean accuracy of the model: 0.38461538461538464
```

```
Parameters used: ('liblinear', True, 'l2', 'balanced')
Mean accuracy of the model: 0.46153846153846156
```

```
Parameters used: ('liblinear', False, 'l1', None)
Mean accuracy of the model: 0.5384615384615384
```

```
Parameters used: ('liblinear', False, 'l1', 'balanced')
Mean accuracy of the model: 0.46153846153846156
```

```
Parameters used: ('liblinear', False, 'l2', None)
Mean accuracy of the model: 0.3076923076923077
```

```
Parameters used: ('liblinear', False, 'l2', 'balanced')
Mean accuracy of the model: 0.46153846153846156
```

Note that in the graded assignment, you will take a more generalizable approach that doesn't require you to explicitly specify each hyperparameter.

That is, instead of:

```
LogisticRegression(solver=hp[0], fit_intercept=hp[1], ...
```

You'll be able to write:

```
LogisticRegression(**params)
```

Looks like none of these models beats the original! This won't always be the case, so next time the opportunity arises, you'll be able to check this for yourself.

Grid Search

This is essentially grid search. You'll be implementing customized grid search in the graded assignment.

- Note that even though sci-kit learn provides a grid search function, it uses K-fold cross validation, which you won't want to do in the assignment, which is why you will implement grid search yourself.

Congratulations on completing this lecture notebook!

By now, you should feel more comfortable with the `sklearn` library and how it works. You also created a grid search from scratch by leveraging the `itertools` library. Nice work!